



BrainRouter

THE MEMORY LAYER FOR AI AGENTS

Agents that remember.

Capture what matters. Forget what doesn't. Surface the right context on every prompt.

[Live on npm](#)

[Open source](#)

Anh Dang · [Event]



● THE PROBLEM

Every session, your agent starts from zero.

LLMs have no memory between conversations. Your agent re-learns the same facts, re-reads the same files, and re-asks the same questions — every time.

The smartest model in the world **still has amnesia.**



First chat

context rebuilt from scratch



A week later

context rebuilt from scratch



A month later

context rebuilt from scratch

● WHY IT'S HARD

Today's fixes don't scale



Dump the history

Stuff every message into the prompt — it blows the context window and burns tokens.



Flat vector DB

Returns what's mathematically close, not what's useful. Noise in, noise out.



Static prompts

Hand-written rules with no feedback loop. Nothing improves; the agent never learns.

The result: expensive, forgetful, and never improving.

● THE INSIGHT

Model memory like the brain

Short-term feeds long-term. Unused facts fade. Used ones get reinforced.



● WHAT WE BUILT

A cognitive memory engine

Four layers, each with its own lifetime — short-term feeds long-term; the rest persist or evict based on activity.



SensoryStream

Raw dialogue buffer

Transient



CognitiveRecord

Classified facts that decay

Long-term



ContextualFocus

Heat-scored task scenes

Medium




CoreIdentity

Profile + hard rules

Permanent

● HOW RECALL WORKS

The right memories, on every prompt



Three retrievers fused. A judge filters. The graph connects.

Only what's relevant reaches the prompt — a sliver of the context, never the whole history.

- WHY IT'S DIFFERENT

It doesn't just store — it learns



It forgets

Memories fade on a half-life. Instructions stay forever; code facts fade faster than who you are.



It reinforces

When the agent uses a memory, that memory gets stronger — and its decay clock resets.



It judges

An LLM relevance check drops keyword-matches that aren't actually on topic.



It connects

A graph walk surfaces related facts plain keyword or vector search would never reach.

● THE PAYOFF

What that buys you



Fewer tokens

A sliver of the context — not the whole history every time.



Faster replies

Less for the model to read means quicker answers.



Sharper recall

The right memory surfaces — nearly every time.



Flat cost at scale

Spend stays steady as the agent remembers more.

Benchmarked, reproducible, and open — the proof is ready when you are.

● THE PRODUCT

One memory store. Four ways to plug in.



MCP Server

Live

Drop-in memory & skills for any MCP client — Claude Desktop, Cursor, your own agent.



Terminal CLI

Live

A memory-native coding agent with multi-agent orchestration.



Desktop App

Alpha

Native Mac & Windows shell over the same engine.



Web Dashboard

Shipping

See the memory — recalls, contradictions, hosted chat.

All four share the same brain — sign in once, every surface sees the same memory.

● TRACTION

Shipping, in the open



Live on npm

the engine and CLI are public and installable today



Weekly releases

a steady cadence — every release adds capability



Public benchmarks

committed result sets, reproducible by anyone



MIT licensed

open source, end to end

THE PATH SO FAR



● WHY NOW

Memory is becoming its own layer



Agents are everywhere

Coding, support, ops, research — every agent hits the same wall: it can't remember a thing.



A foundational layer

The database was to apps what memory is to agents — infrastructure, not a feature.



Cross-vendor by design

Federation lets agents on any vendor share one brain — the neutral standard.

We're not building a better chatbot. **We're building the memory infrastructure underneath all of them.**

● ROADMAP

Where it's going

● NOW

Memory accuracy, a multi-agent CLI, and a live view of the running fleet.

○ NEXT

Full agent-runtime parity, and the desktop app in alpha.

○ SOON

A plugin & skill marketplace, plus live sync with Git and GitHub.

○ VISION

The memory standard every agent — on any vendor — plugs into.



BrainRouter

Agents that remember.

BrainRouter — the cognitive memory layer for AI agents.



`github.com/kingradiollc/BrainRouter`



`npm i -g @kings/brainrouter-cli`

Let's talk. Anh Dang · anhdang@kingradio.com